APLYD

Governing AI in Oversight Institutions

PRINTABLES PACK

Cohort 3 · Oversight Institutions

Everything to print and set out before participants arrive: worksheets, case cards, templates, and table tents.

PRINT CHECKLIST — PER PARTICIPANT UNLESS NOTED

1 × Facilitator pre-arrival checklist (you only — next page)
1 × Mental Models reference card
1 × FIT Test worksheet
1 × Build / Buy / Partner decision card
6 × Forensic case cards (1 set per table)
1 × Vendor seven-probe card (1 per trio)
1 × Crisis tabletop pack (1 per table) + 1 comms template + 1 AAR sheet each
1 × AI Committee charter template + 1 × Model Card template (per table)
1 × First-90-days plan + 1 × Letter to Self + envelope
1 × Table tent (name card)

Facilitator pre-arrival checklist

Trainer only — work through this the day before.

Room

- Tables of 4–6, cabaret style (not classroom rows).
- Projector + the C1 deck loaded; clicker tested.
- Flip charts + markers at each table.
- Wall space for the opening-poll capture.

Print & place

- Set the Mental Models card, FIT worksheet, and Build/Buy/Partner card at each seat.
- Place one set of the six case cards per table.
- Place the crisis pack face-down — released only at the tabletop.
- Letter to Self template + a sealable envelope per seat (collect at close).

Confirm before Day 1

- The faculty roster is inserted into the deck (replace placeholder cards).
- The OPEN governance facts (institutional protections, escalation) are confirmed for your context.
- Each participant has been asked to bring ONE real deployment from their portfolio.
- Pre-reads (Participant Handbook §2) sent at least one week ahead.

Mental Models — reference card

Eight lenses for reading any AI deployment. Keep this in front of you all programme.

#	Mental model	The question to ask
1	Accuracy is the most lied-about number	Measured how, on whom, by what definition?
2	AI works on average; your job is the tails	Which subgroup does it fail?
3	AI failure is organisational	Wrong incentive / authority / appeal — fix that first.
4	AI is 10% of the work	Who owns the other 90%?
5	Models reconstruct attributes from proxies	What proxy hides the bias?
6	Failure scales faster than correction	Match correction speed to failure speed.
7	The override's design decides if it works	Has it been tested under pressure?
8	A policy can be encoded without being legislated	Would it survive being written down?

FIT Test — worksheet

Score your deployment. One Critical is a stop condition — do not average.

Deployment:

FIT question	P / W / C	One-line defence
F1 · Does the tool's one job match the decision you face?		
F2 · Is it the right tool for THIS population?		
I1 · What is the subgroup performance?		
I2 · What data was it trained on?		
I3 · What is the failure mode — and who does it fall on?		
T1 · Can you monitor, override, appeal, and sunset it?		
T2 · Who is accountable — and can they pause within 24h?		

Verdict (Pass / Watch / Critical) and why:

Build / Buy / Partner — decision card

Door	Choose it only when…
BUILD	It is core for 5+ years AND you can hire and keep the talent AND you can carry the lifecycle (data, monitoring, sunset).
BUY	It is a commodity AND the vendor has field evidence in your context AND the exit/continuity clause is signed up front.
PARTNER	No credible vendor exists yet AND the partner fills a specific named gap AND joint governance and sunset are agreed.

My capability:

My choice + 250-word defence against the other two paths:

Case A · Toeslagenaffaire (Netherlands)

Failure forensics — discussion card

What happened. A tax-authority risk model flagged childcare-benefit fraud, scoring foreign-sounding names and dual nationality higher. ~26,000 families wrongly accused; ~2,090 children removed; the cabinet fell in 2021.

The lesson. The algorithm was a contributor; the cause was binding automated decisions with no proportionate appeal.

> **WHERE COULD IT HAVE BEEN CAUGHT?**
> An independent appeal pathway and a discrimination audit of features — at design, not after the harm.

Source: Dutch Parliamentary Inquiry, 2020; Amnesty International, 2021.

Case B · Robodebt (Australia)

Failure forensics — discussion card

What happened. Automated income-averaging raised ~470,000 wrongful debts; ruled unlawful; Royal Commission 2023; A$1.8bn in settlements.

The lesson. Harm scaled at machine speed; correction ran at the speed of courts.

> **WHERE COULD IT HAVE BEEN CAUGHT?**
> A pre-launch legal review of the averaging assumption, and an appeal system sized to the rate of debt-raising.

Source: Royal Commission into the Robodebt Scheme, 2023.

Case C · Kenya digital-lending collapse (Kenya)

Failure forensics — discussion card

What happened. From 2016, 800+ unlicensed apps issued instant micro-loans, harvested borrowers' phone contacts, and shamed defaulters by calling employers, relatives and friends; 8M+ Kenyans were blacklisted. In 2020 the regulator revealed a man took his life after such harassment.

The lesson. Automated lending plus harvested data scaled harassment and exclusion to millions before any regulator had authority; harm far outran correction.

> **WHERE COULD IT HAVE BEEN CAUGHT?**
> A licensing perimeter before launch, a ban on third-party debt-shaming, and a proportionate default threshold for credit-bureau listing.

Source: Central Bank of Kenya, 2020–22; Daily Nation, 2020; CBK Digital Credit Providers Regulations 2022.

Case D · COMPAS + predictive policing (US)

Failure forensics — discussion card

What happened. A court risk tool showed ~2× the false-positive rate for Black defendants; predictive policing entrenched over-policing. Challenges largely failed.

The lesson. A model trained on biased history launders that bias into 'evidence.'

> **WHERE COULD IT HAVE BEEN CAUGHT?**
> A disparate-impact audit before deployment, and a ban on using arrest data as a proxy for crime.

Source: ProPublica, 2016; Loomis v. Wisconsin, 2016.

Case E · Aadhaar exclusion (India)

Failure forensics — discussion card

What happened. Biometric authentication for rations excluded the elderly, manual labourers, and low-connectivity areas. ~64M cards cancelled; 92% success still excludes millions; starvation deaths documented.

The lesson. Aggregate success rates hide systemic exclusion; no fallback means exclusion by design.

> **WHERE COULD IT HAVE BEEN CAUGHT?**
> A mandatory non-digital fallback, and reporting the count excluded — not the percentage who succeeded.

Source: Puttaswamy, 2018; CAG PDS reports 2019–22.

Case F · Colombia's 'secret algorithm' (Colombia)

Failure forensics — discussion card

What happened. Facing COVID, Colombia sent cash to ~3 million households chosen by the Sisbén database plus a confidential algorithm scoring 'income-generating capacity.' The government refused to disclose how it worked; excluded families could not see why, or appeal.

The lesson. A secret, unauditable eligibility algorithm has no accountability — even when it reaches millions fast.

> **WHERE COULD IT HAVE BEEN CAUGHT?**
> Public criteria and a contestation pathway designed in from the start; transparency is the legitimacy of the system, not an optional extra.

Source: CHRGJ, 'Chosen by a Secret Algorithm,' 2021; Privacy International, 2022.

Vendor diligence — the seven probes

Read these aloud in the meeting. Stay on each until it is answered.

1. Show me sensitivity — not accuracy — on a population like ours.
2. Show me subgroup performance, broken out, in writing.
3. What data was this trained on, and how does it differ from our context?
4. What are the known failure modes, and who do they fall on?
5. What is the exit / continuity plan if you are acquired or fail?
6. Who owns the model, the data, and the outputs?
7. What does it cost over five years — integration, monitoring, retraining, sunset — not just licence?

The probe the vendor could NOT answer (your diligence finding):

Crisis tabletop — the Day-47 scenario

Released by the facilitator at the tabletop. The room is the crisis cell.

Inject 1. Day 47 of your deployment. A journalist emails: families in one district report being wrongly cut off by your system. They want comment in 3 hours.

Inject 2 (release after first decisions). Your data team confirms a subgroup the model mis-scores. ~4,000 people affected so far.

Inject 3. A political principal calls: “Is this our fault? What do I say?”

Your tasks: (1) Contain — pause? who authorises, how fast? (2) Communicate — draft the holding statement. (3) Correct — how do you reprocess the affected cases? Then run the After-Action Review.

Crisis — holding statement template

What we know (plain, honest, no blame):

What we are doing right now:

What affected people should do:

When we will update next:

Crisis — After-Action Review

What happened (timeline):

What worked:

What failed — and was it design, deployment, or appeal?

The one control we will add before we restart:

AI Committee — charter template

Purpose (what this committee governs):

Membership (names / roles — include an independent voice):

Authority (what it can require, approve, and PAUSE):

What it must see for every deployment (model card, subgroup data, appeals log):

Cadence (how often it meets — and the emergency trigger):

Named pause authority + maximum time to pause:

Model Card — template

Deployment name / version:

What it does (one job):

On whom (population) and the decision it informs:

Training data — source, date, and how it differs from our context:

Performance — overall AND by subgroup:

Known failure modes — and who they fall on:

Human override — how it works and when it was last tested:

Sunset conditions (when we pause or retire it):

Accountable owner:

RACI — template for an AI decision

For one AI decision in your deployment, name who is Responsible, Accountable, Consulted, Informed. Exactly one Accountable per row.

Decision / step	Responsible	Accountable	Consulted	Informed
Approve the deployment				
Validate model performance + subgroups				
Authorise a pause / kill switch				
Handle an appeal / contested decision				
Sign off public-facing AI statements				
Decide sunset / rebuild				

Sunset checklist

Complete before launch, not at the end. A deployment with no sunset plan is one you cannot stop.

Sunset triggers (performance, harm, cost, legal, or end-of-need conditions that pause or retire it):

Who decides to pause, and within how long:

What continues for the people it served when it stops (the non-AI fallback):

Data: what is retained, deleted, or returned on sunset:

Vendor exit / continuity terms (signed at procurement):

Date this plan was last reviewed:

First 90 days — plan

Three governance actions, named owners, dates.

Action	Owner	By when

The single deployment I will govern first, and why:

Letter to Self

Sealed at the close. Returned to you at the 90-day cohort call. Private — no one else reads it.

In the next 90 days, I commit to do the following differently:

Signed: _____ Date: _____

Feedback & evaluation

One per participant, collected at close. Anonymous.

Cohort / date:

Rate 1 (low) to 5 (high):

	1	2	3	4	5
The programme improved my ability to govern AI					
The cases were relevant to my work					
The activities were useful					
The materials (handbook, printables) were useful					
I would recommend this to a colleague					

The single most useful thing:

The single thing to improve:

One AI deployment I will act on differently because of this:

Certificate of completion

APLYD

Governing AI in Oversight Institutions

This certifies that

---

has completed the APLYD Cohort 3 · Oversight Institutions programme,

demonstrating applied competence in interrogating, procuring, governing, and sustaining AI deployments in the public interest.

Programme lead: _____ Date: _____

Table tent (name card)

Print, fold along the centre, stand on the table. One per participant.

[NAME]

— — — — — — — — — — fold here — — — — — — — — — —

APLYD · Governing AI in Oversight Institutions